

Agentic Solution Opportunities Across Accenture's Client Base

I'll be upfront about what I know and don't know: I have general knowledge about Accenture's industry coverage and the kinds of problems their clients typically face. I don't have insider access to specific account plans or opportunities. But Accenture publishes a lot about their industry practices, and the patterns of what's sellable are fairly consistent. So take this as **an informed map of the terrain**, not insider intel.

The goal: help you see where agentic solutions genuinely fit, beyond copilots. Because "agentic" is much bigger than "chatbot with tools."

First — Reframe What "Agentic" Means Commercially

Most people equate agentic with "AI chatbot that uses tools." That's one shape. The actual pattern is broader:

Agentic = software that pursues a goal across multiple steps, uses tools/systems autonomously, reasons about what to do next, and handles branching logic that's too complex for deterministic workflows.

Once you see it this way, agentic solutions fit **anywhere there's an expert human doing a multi-step investigative or judgment-heavy task inside a tool-rich environment**. That's a huge surface area.

Three commercial shapes to keep in mind:

1. **Copilot** — human stays in the loop, agent augments. (Lowest risk, fastest sale.)
2. **Autopilot with approval gates** — agent executes, human approves at key moments. (Medium risk, high ROI.)
3. **Full autonomy** — agent runs end-to-end, only escalates exceptions. (Highest risk, highest value.)

Most enterprise sales land at #1 or #2. #3 is the 3-year roadmap.

Industry-by-Industry Agentic Opportunities

Financial Services

Accenture's largest vertical. Enormous appetite for AI, big compliance constraints. Agentic fits beautifully because **every claim needs a citation and every action needs an audit trail** — which you've already designed for.

Opportunity	Who in the org buys it	Shape
KYC / onboarding investigation agent — ingests documents, cross-references sanctions lists, company registries, adverse media; drafts analyst memo with citations	AML / Financial Crimes	Copilot → Autopilot
Alert triage for transaction monitoring — closes false positives, escalates real suspicion with evidence package	BSA/AML operations	Autopilot with approvals
Credit underwriting assistant — pulls financials, runs models, drafts credit memo	Commercial banking	Copilot
Dispute resolution agent — investigates chargebacks across systems, proposes resolution	Card operations	Autopilot
Wealth advisor research agent — pulls portfolio data, market news, regulatory constraints, proposes rebalance	RIA / Private banking	Copilot (always — regulated)
Insurance claims adjudication agent — ingests claim, pulls policy, damage photos, medical codes; proposes settlement	P&C insurance	Autopilot with approvals
Trade break investigation agent — reconciles across systems, proposes fix	Capital markets ops	Autopilot

Why it sells: These are jobs where senior analysts spend hours on evidence-gathering and junior analysts need supervision. The ROI math is straightforward — FTE cost reduction or throughput multiplier. Compliance loves the citation pattern.

Healthcare & Life Sciences

Huge Accenture vertical, especially in payers, providers, and pharma. Enormous text-heavy workflows.

Opportunity	Who buys	Shape
Prior authorization agent — ingests request, pulls medical records, policy rules, prior denials; drafts decision	Health payers (UnitedHealth, Anthem, etc.)	Autopilot with MD approval
Medical coding agent — reads clinical notes, proposes CPT/ICD codes with citations	Providers, RCM companies	Copilot
Clinical trial matching agent — reads patient record, trial protocols; proposes eligible trials	Pharma, AMCs	Copilot
Regulatory submission assembly agent — pulls study data, drafts modules of submission	Pharma regulatory affairs	Copilot
Pharmacovigilance case processing agent — ingests adverse event reports, codes them, drafts narrative	Pharma safety	Autopilot with approvals
Provider network adequacy agent — checks network against regulatory rules, suggests gaps	Payers	Autopilot
Medical literature review agent — for MSLS (medical science liaisons)	Pharma commercial	Copilot

Why it sells: Prior auth alone is a \$20B+ problem. Medical coding is another multi-billion-dollar market. These are unglamorous but massive.

■ Energy & Utilities

Accenture's Industry X practice is strong here. Asset-heavy, data-rich, increasingly tech-driven.

Opportunity	Who buys	Shape
Outage diagnosis & crew dispatch agent — correlates SCADA, weather, AMI data; proposes crew routing	Distribution utilities	Autopilot
Asset inspection report agent — ingests drone imagery + sensor data; drafts inspection findings	T&D utilities, oil & gas	Copilot
Regulatory filing preparation agent — pulls operational data, drafts FERC/rate case filings	Utility regulatory teams	Copilot
Trading floor assistant — news, weather, load forecasts, regulations affecting trades	Power/gas traders	Copilot
Well performance investigation agent — anomaly detection + diagnosis across wells	Oil & gas upstream	Copilot → Autopilot
Vegetation management agent — satellite imagery + LiDAR → work orders	Utilities	Autopilot

Why it sells: Regulatory pressure, aging workforce, massive asset bases. “We can augment your senior engineers” is a winning pitch.

Manufacturing & Supply Chain

Opportunity	Who buys	Shape
Supplier risk investigation agent — financial health, geopolitical exposure, ESG, news monitoring	Procurement / supply chain	Copilot
Supply chain disruption response agent — detects disruption, proposes rerouting, alternate suppliers	Control tower ops	Autopilot with approvals
Quality incident RCA agent — correlates manufacturing data, proposes root cause	QA engineering	Copilot
Warranty claims investigation agent — proposes decision based on records	Automotive, appliances	Autopilot
Engineering change request agent — impact analysis across BOMs, suppliers, compliance	PLM teams	Copilot
Shop floor troubleshooting copilot — reads MES + PLC data, proposes fix	Plant engineering	Copilot

Why it sells: Shortage of skilled workers, supply chain fragility post-COVID, massive institutional knowledge leaving with boomer retirements.

Retail & Consumer Goods

Opportunity	Who buys	Shape
Merchandising assistant — reads sales data, competitor pricing, inventory; proposes assortment/pricing	Buyers/planners	Copilot
Trade promotion planning agent — proposes promo calendars based on history + forecast	CPG commercial	Copilot
Customer service ops agent — like your telco copilot but for retail (order issues, returns, shipping)	Retail contact centers	Copilot
Store operations agent — reads store KPIs, proposes interventions, drafts communications	District managers	Copilot
Content & PDP generation agent — product descriptions, SEO, A/B variants	E-commerce merch	Autopilot
Shrinkage / fraud investigation agent — pattern detection across transactions, CCTV metadata	Loss prevention	Copilot

Travel, Transportation & Hospitality

Opportunity	Who buys	Shape
Irregular operations (IROPS) agent — when flights cancel, re-accommodate pax with best options	Airline ops	Autopilot with approvals
Revenue management assistant — pricing decisions with market context	Airlines, hotels	Copilot
Customer compensation adjudication — claim eligibility, regulatory rules (EU261), goodwill limits	Airline customer care	Autopilot
Fleet maintenance diagnosis agent — reads telematics, MEL rules; proposes action	Airlines, rail, trucking	Copilot
Concierge agent for premium guests — researches, books, personalizes	Luxury hospitality	Copilot

Why it sells: Airlines especially have relentless cost pressure and IROPS alone destroys NPS — this is a high-impact wedge.

Public Sector

Accenture Federal Services and equivalent public sector practices. High scrutiny, high volume, low-budget environments — but increasingly open to AI.

Opportunity	Who buys	Shape
Benefits eligibility investigation agent — reads applications, policy rules, prior decisions	State/federal benefits agencies	Copilot
Tax examination assistant — reads returns, flags issues, drafts notices	Revenue agencies	Copilot
FOIA / public records processing agent — redaction, routing, response drafting	Federal/state records teams	Autopilot with approvals
Case management copilot for social workers	HHS, state agencies	Copilot
Procurement / RFP response agent — for contractors responding to gov RFPs	Defense/civilian contractors	Copilot

Compliance note: FedRAMP, CJIS, ITAR apply. Bedrock has gov-ready variants which helps.

Legal, Tax, Audit (Professional Services)

A surprising sweet spot. Accenture sells to law firms and Big 4 clients, plus corporate legal/tax/audit functions.

Opportunity	Who buys	Shape
Contract review agent — redlines against playbook, flags risks, proposes alternates	Corporate legal	Copilot
Tax research agent — multi-jurisdictional analysis with citations	Corporate tax, Big 4	Copilot (always cited)
Audit workpaper preparation agent — pulls evidence, drafts testing memos	Internal audit, external audit	Copilot
Due diligence agent — M&A investigation across data rooms	Corp dev, PE firms	Copilot
Litigation document review agent — already exists, but agentic versions are the next gen	Law firms, corporate litigation	Autopilot with review

Why it sells: Billable hour pressure, associate attrition, AI-native competitors (Harvey, Hebbia, Robin) forcing incumbents to act.

Real Estate & Construction

Less AI-saturated market = more greenfield.

Opportunity	Who buys	Shape
Lease abstraction + management agent	CRE owners, brokers	Autopilot
Construction RFI/submittal agent	GCs, project management	Copilot
Permit/zoning research agent	Developers	Copilot
Property underwriting agent	REITs, PE real estate	Copilot

Higher Education

Smaller deals but real opportunities.

Opportunity	Who buys
Student advising agent — degree audit, course recommendations, policy	Registrar, advising
Admissions triage agent	Enrollment management
Research grant application agent	Sponsored research offices

Internal (Enterprise-Wide) Agentic Opportunities

Across all verticals, these horizontal plays sell everywhere:

Opportunity	Who buys	Shape
SecOps Tier 1 analyst copilot — triages alerts, enriches, proposes disposition	CISO org	Copilot → Autopilot
DevOps incident response agent — reads telemetry, logs, runbooks; proposes fix	SRE teams	Copilot
ITSM / Service desk agent (internal employee support)	CIO org	Copilot/Autopilot
HR employee query agent — benefits, policy, leave, payroll	HR ops	Autopilot
Procurement / P2P agent — PO questions, invoice matching, approvals	Finance ops	Autopilot
Sales research agent — account planning, prospect research, prep briefs	Enterprise sales	Copilot
Deal desk / pricing approval agent	Revenue ops	Copilot
Finance close copilot — reconciliations, flux analysis, draft narratives	Corporate finance	Copilot
Internal audit testing agent	Internal audit	Copilot
Data engineering / lineage investigation agent	Data platform teams	Copilot

These are often the easiest first sale into any enterprise because the buyer is internal IT/ops, not a risk-averse line of business.

What Makes an Opportunity Actually Sellable (Filter)

Use this filter on any idea:

1. **Is there a senior human doing repetitive investigation today?** If yes → candidate.
2. **Are the tools/data digital and API-accessible?** If no → not yet.
3. **Are outcomes measurable?** (AHT, throughput, error rate, hours saved.) If no → hard to sell.
4. **Are mistakes recoverable or gated?** If no → needs heavy compliance work.
5. **Is there existing organizational urgency?** (Budget, mandate, fear.) If no → education sale, longer cycle.
6. **Can you show value in 6–12 weeks?** If no → doesn't fit Accenture's pilot model.

The best opportunities score yes on all six. Your telco copilot does.

The Meta-Play: Build a Portfolio, Not a Product

Here's the strategic insight most people miss:

You shouldn't pick *one* vertical and build deeper. You should build a **portfolio of 3–4 reference copilots across different industries** that all share the same underlying agentic architecture.

Why:

1. You demonstrate **architectural generality** — “same core, different tools” is the most compelling enterprise AI story right now.
2. You become the **go-to agentic expert in your practice** across deals, not just in one vertical.
3. Each reference asset costs you maybe 2–3 weeks to produce once you have the core. ROI is enormous.
4. Accenture partners in different verticals will pull you into their deals.

A concrete portfolio I'd build in your position:

1. **Telco support copilot** ✅ (done)
2. **SecOps Tier 1 analyst copilot** (horizontal, sells everywhere, high demand)
3. **Financial crime / KYC investigation copilot** (banking, regulated, high-value)
4. **Healthcare prior auth copilot** (massive market, distinctive)

Four assets, shared core, covers maybe 60% of Accenture's biggest verticals.

One Honest Warning

Every Big 4 consultancy and every systems integrator is pitching “agentic AI” right now. **Differentiation is not the agent — everyone has one.** Differentiation is:

1. **Domain fidelity** — does it actually reflect how the work is done?
2. **Integration depth** — does it plug into real client systems, not mocks?
3. **Risk handling** — approval gates, citations, audit trails, graceful degradation.
4. **Evaluation rigor** — can you prove it works on real data?
5. **Deployment muscle** — can Accenture actually ship it at scale?

Your POC nails #3. Partially nails #1. Needs work on #2, #4, #5 — which is exactly what a client engagement pays for.

Don't compete on “we have agents.” Compete on “we have agents that survive your compliance team, your ops team, and your CIO.” That's a much shorter list.

TL;DR

- Agentic solutions fit virtually every Accenture client because every enterprise has senior humans doing multi-step, tool-heavy investigation work.
- The highest-value verticals are **financial services, healthcare, energy, and enterprise-internal (SecOps/DevOps/ITSM)**.
- Build a **portfolio of 3–4 reference assets** sharing one core architecture, not a single product.
- Differentiate on **risk, audit, integration, and evaluation** — not on “we built an agent.”
- Your telco copilot is the right first asset. The second one should be horizontal (SecOps or ITSM) to maximize reach inside Accenture.

Want me to go deep on the architecture for any specific one of these — e.g., how a SecOps copilot or a KYC investigation agent would differ structurally from your telco copilot?

Top 10 Most Sellable Agentic POCs — Ranked by Commercial Viability

I’ve optimized this list on three axes simultaneously:

1. **Novelty** — not another “chatbot for X.” Either solves a problem clients don’t know they have, or solves a known problem in a way the market hasn’t productized yet.
2. **Feasibility** — you can build a credible POC in 2–4 weeks with your current stack (Bedrock + Python + the architecture you’ve designed).
3. **Sellability** — there’s an identifiable buyer with budget, a measurable ROI story, and a deal size that justifies Accenture’s involvement (\$500k+ pilots).

I’ve deliberately avoided the obvious ones (generic “customer support copilot,” “HR chatbot,” “sales email generator”) because those are commoditized. Each of these has a real wedge.

#1 — The Controls & Compliance Evidence Agent

Industry: Any regulated enterprise (banks, insurers, healthcare, pharma, utilities) **Novelty:** 8/10 • **Feasibility:** 9/10 • **Sellability:** 10/10

The problem they don’t think they have

Every enterprise runs dozens of compliance control programs — SOX, SOC 2, HIPAA, PCI, ISO 27001, model risk (SR 11-7), GDPR, etc. **Evidence collection** consumes armies of GRC analysts who spend 60% of their time screenshotting, exporting logs, writing “tested and passed” memos, and routing them to auditors.

Nobody has framed this as an agentic problem. Everyone treats it as a GRC platform problem (Archer, ServiceNow GRC) — but those platforms are workflow tools, not intelligent collectors.

What the agent does

Given a control ("All production database access is logged and reviewed monthly"), the agent:

1. Identifies the systems to query (via a control-to-system mapping).
2. Calls APIs to pull evidence (AWS CloudTrail, Okta, Splunk, database audit logs).
3. Validates completeness against the control language.
4. Drafts the testing memo with citations.
5. Routes to human for approval.

Why it sells

- **Every F500 has this pain.** Universal.
- **ROI is concrete:** 40–70% reduction in evidence-collection hours.
- **Audit-defensible by design** — citations + trace logs are the point, not a feature.
- **CISO, Chief Compliance Officer, CFO** all want this. Three buyers.
- **Repeatable across every client.** Near-zero customization for the core; just add connectors.

Why it's feasible for you

Your citation enforcement + risk tiers + observability design is **purpose-built for this use case**. You'd build it faster than anyone else approaching it fresh.

#2 — The Alert Fatigue Triage Agent (SecOps / Tier-1 SOC)

Industry: Any enterprise with a SOC — financial services, healthcare, critical infrastructure **Novelty:** 7/10 • **Feasibility:** 9/10 • **Sellability:** 10/10

The problem

SOC analysts drown in 5,000–50,000 alerts/day. 95%+ are false positives. Analysts burn out, real threats get missed, and MDR vendors charge a fortune for what is effectively triage labor.

"AI SOC" startups (Dropzone, Prophet, Radiant) are proving this market hard — which is exactly why Accenture clients want to consider it without buying a startup they can't trust yet.

What the agent does

For each alert:

1. Enriches from SIEM, EDR, threat intel, identity provider, asset DB.
2. Runs an investigation plan (same pattern as your telco copilot — hypothesis → tool calls → evidence).
3. Classifies: false positive (auto-close with evidence), benign (deprioritize), suspicious (escalate with pre-built case).
4. Writes the investigation narrative analysts would have written.

Why it sells

- **Board-level problem.** CISOs are personally accountable.
- **Huge budget.** Security teams spend \$10M+/year on MDR services; this displaces a chunk.
- **Measurable in weeks:** alerts closed per hour, MTTR, analyst hours recovered.
- **Accenture has deep SecOps practices** (Accenture Security is a multi-billion-dollar business).
- **Every industry buys it.** Horizontal play.

Why it's feasible

Your architecture is literally shaped like a SOC investigator: evidence log + multi-step diagnosis + citation + approval for containment actions. Swap tools, repurpose the loop.

The strategic kicker

This is the **second asset I'd build after telco**. Maximum reach inside Accenture — every industry team pulls you into their deals.

#3 — The Shadow IT & SaaS Sprawl Investigator

Industry: Horizontal (every CIO/CISO has this; enterprise-wide buyer) **Novelty:** 9/10 • **Feasibility:** 8/10 • **Sellability:** 8/10

The problem they don't know they have

Every F500 has 400–1,200 SaaS apps in use. IT knows about maybe 150 of them. The rest are bought on corporate cards by individual teams. This creates:

- Security exposure (unknown data leaves the org)
- Massive spend duplication (12 different note-taking apps)
- Compliance risk (HIPAA data in unapproved tools)
- Vendor management chaos

SaaS management platforms (Zylo, Productiv, Torii) find the apps. **Nobody uses agents to investigate and act on them.**

What the agent does

1. Ingests SSO logs, expense reports, DNS traffic, browser extension data.
2. For each unknown app: researches what it is, who uses it, what data it touches.
3. Cross-references against approved vendor list, contracts, security reviews.
4. Drafts a disposition: consolidate, approve, block, renegotiate.
5. For high-risk: opens a ticket to the right owner with evidence.

Why it sells

- **CFO cares** (spend waste). **CISO cares** (data exfil). **CIO cares** (architecture sanity).
- **Hard-dollar savings story** — typical F500 can reclaim \$5M–\$30M/year on SaaS spend alone.
- **Very low competition** — the category is “SaaS management,” but nobody is running agents over the data.

Why it's feasible

Tools are all API-accessible (Okta, Expensify, NetSkope, DNS logs). Mock integrations work fine for the demo.

The novelty angle

This is a problem every client **has but hasn't named**. Walking into a CIO's office with “did you know you have 847 SaaS apps and here's what 312 of them do” is a conversation-stopper. That's what makes it sellable.

#4 — The “Deal Risk” Investigation Agent for M&A Due Diligence

Industry: Private equity, corporate M&A, investment banks **Novelty:** 9/10 • **Feasibility:** 7/10 • **Sellability:** 9/10

The problem

Due diligence for M&A is brutally manual: data rooms with 10,000 documents, compressed timelines (often 2–4 weeks), and junior analysts reading contracts at 2am looking for change-of-control clauses, MAC provisions, indemnity carve-outs, or hidden liabilities.

Existing tools (Luminance, Kira, Harvey) do contract analysis. **Nobody builds an investigative agent that works across contracts, financials, customer lists, and public data simultaneously to build a risk profile.**

What the agent does

Given a target company's data room + a deal thesis:

1. Classifies documents (contracts, financials, HR, IP, litigation).
2. Investigates specific risk questions: "What % of revenue depends on contracts with change-of-control triggers?" "Are there undisclosed related-party transactions?" "What's the customer concentration risk?"
3. Cross-references findings with external sources (litigation databases, news, regulatory filings).
4. Builds a risk memo with citations to exact contract paragraphs.

Why it sells

- **PE firms are printing money and have AI budgets.** KKR, Blackstone, Apollo all have internal AI teams.
- **\$1M+ engagements routinely.** This is partner-level work.
- **Accenture has a huge M&A advisory practice** (financial advisory services).
- **Dramatic time savings** — 4-week DD becomes 2-week DD. That's worth real money to dealmakers.

Why it's feasible (mostly)

Document handling is harder than your telco POC (RAG over a data room). But the agentic loop is the same. Stretch: 4–6 weeks instead of 2–3.

Why the novelty score is high

Most legal AI is **document-by-document analysis**. This is **thesis-driven investigation**. Different paradigm.

#5 — The Pricing & Discount Anomaly Agent (Revenue Leakage)

Industry: B2B SaaS, manufacturing, distribution, telecom **Novelty:** 8/10 • **Feasibility:** 9/10 • **Sellability:** 9/10

The problem they don't think they have

Every B2B company with a sales force and a price list has **revenue leakage**: unauthorized discounts, incorrect quotes, wrong entitlements, missed price increases, contracts where the customer is getting terms better than the current book price. Deloitte estimates 1–3% of B2B revenue leaks this way. For a \$1B company, that's \$10–30M/year.

Finance knows this happens but can't investigate systematically.

What the agent does

1. Continuously scans CPQ, ERP, CLM, and order systems.
2. Investigates anomalies: "This renewal priced 40% below peer accounts — why?" "This customer has a discount tier they shouldn't qualify for."
3. Reviews the contract to validate.
4. Drafts a recovery memo: "Propose price increase at next renewal, estimated \$420K recovery."
5. Routes to account team for approval/action.

Why it sells

- **Self-funding ROI.** If the POC finds \$5M of leakage at one client, it pays for itself 10x over.
- **CFO buyer.** Clear line of sight to P&L impact.
- **Deloitte, EY, PwC already sell "revenue leakage consulting"** as services. You productize the core analytical work.
- **Every industry with contracted pricing applies.** Massive TAM.

Why it's feasible

All data sources are structured (Salesforce CPQ, SAP, Zuora, etc.). Agent reasoning over structured data is your wheelhouse.

The sales pitch

"We'll run this for 4 weeks in read-only mode. We only get paid if we find leakage." Outcome-based pricing. Irresistible.

#6 — The Regulatory Change Impact Agent

Industry: Banking, insurance, pharma, medical devices, energy **Novelty:** 9/10 • **Feasibility:** 8/10 • **Sellability:** 9/10

The problem

Regulators publish dozens of proposed and final rules per year (FINRA, SEC, OCC, FDA, FERC, DORA, EU AI Act, etc.). Each rule change triggers an internal assessment: which policies, controls, systems, and products are affected? Today this is a weeks-long exercise involving legal, compliance, risk, and product teams.

Companies like Ascent and [Compliance.ai](#) do rule ingestion. Nobody does automated impact analysis.

What the agent does

Given a new regulatory rule:

1. Parses the rule into requirements.
2. Searches the internal policy library, control library, system inventory for affected items.
3. Identifies gaps: "Rule requires retention of X for 7 years; current policy says 5."
4. Drafts an impact memo and change-request list for each affected team.
5. Tracks implementation across teams.

Why it sells

- **GRC and regulatory affairs teams are drowning** — especially with the volume of AI/data/crypto/climate rules now appearing.
- **Very visible at the board level.** Fines for missing regulatory deadlines are enormous.
- **Accenture Risk practice already has deep relationships here.**
- **Defensible with audit trails** — critical for regulated industries.

Why it's feasible

Regulatory text + internal docs = RAG. Impact analysis = agentic reasoning. Both play to your stack.

The wedge

Start with **one regulation** (e.g., DORA in EU, or a specific FDA guidance) and one client. Prove the pattern. Every other regulation is a repeat sale.

#7 — The "Silent Failure" Detection Agent for Enterprise Data Pipelines

Industry: Any data-mature enterprise (banking, retail, healthcare, tech) **Novelty:** 10/10 • **Feasibility:** 8/10 • **Sellability:** 8/10

The problem nobody is talking about

Modern data stacks (dbt, Snowflake, Airflow, Databricks) have **observability tools** for pipeline failures. But **silent failures** — where pipelines succeed but produce wrong data — are far more dangerous. Schema drift, upstream changes, logic errors, partial loads. Teams find out weeks later when an executive notices numbers look off.

Observability tools (Monte Carlo, Bigeye, Anomalo) detect statistical anomalies. **They don't investigate root cause.**

What the agent does

When an anomaly is flagged (or proactively, on a schedule):

1. Investigates: "Revenue for region X dropped 30% day-over-day. Why?"
2. Traces lineage upstream.
3. Checks recent schema changes, dbt model changes, source system changes.
4. Runs diagnostic queries.
5. Drafts a root cause memo: "Source system changed currency column from USD to local currency on Tuesday. Affects 14 downstream models."
6. Proposes the fix.

Why it sells

- **Data quality is finally a board-level issue** because it breaks AI initiatives.
- **Chief Data Officer is a real buyer with real budget** now.
- **Very few people are solving this.** Mostly startups.
- **Accenture has a huge data practice.** Natural fit.

Why the novelty is maxed out

The pattern "agent that investigates its own data infrastructure" is genuinely fresh. Clients light up when they see it.

Why it's feasible

Lineage tools, dbt metadata, warehouse query logs — all API-accessible. Agent reasoning over structured metadata is easy.

#8 — The RFP Response Agent (for Companies That Respond to RFPs)

Industry: Defense contractors, B2B software, consulting firms, construction, staffing **Novelty:** 7/10 • **Feasibility:** 9/10 • **Sellability:** 9/10

The problem

Large enterprises respond to 50–500 RFPs/year. Each takes 2–6 weeks and consumes proposal teams, subject matter experts, legal, pricing. Win rates hover around 20–30%. Meanwhile, most answers are 60–80% reusable from prior RFPs that nobody can find.

Tools like Loopio, Responsive exist — they're **databases**, not agents.

What the agent does

Given an incoming RFP:

1. Parses the RFP into discrete questions and requirements.
2. For each question: searches prior RFP responses, marketing content, product docs.
3. Drafts answers with citations.
4. Identifies gaps where SMEs must provide new content.
5. Flags contractual risks that legal must review.
6. Routes sections to the right owners.

Why it sells

- **Sales and proposals leaders have clear ROI.** Win-rate improvement of even 2–3 points is worth millions.
- **Accenture itself responds to RFPs constantly** — you can dogfood it.
- **Every firm with a proposal team is a buyer.**

Why feasibility is high

RAG + agentic composition + human approval gates. Exactly your stack.

Honest caveat

Less novel than others on this list — it's a well-understood category. But the **agentic version** (that actively investigates and composes, vs. just retrieves) is new enough to win deals.

#9 — The Clinical Documentation & Coding Integrity Agent

Industry: Healthcare providers (hospitals, health systems) **Novelty:** 8/10 • **Feasibility:** 7/10 • **Sellability:** 10/10

The problem

Hospitals lose billions annually to **under-coded** claims (services delivered but not billed correctly) and **denied claims** (billing errors caught by payers). Coding companies (3M, Optum) do this with humans + rules-based tools. Startups (Fathom, Nym) do autonomous coding.

The agentic angle no one has cracked: clinical documentation *integrity*. Reading a clinician's note and asking "does this note adequately support the code that was billed?" — and if not, drafting the clarifying query to send back to the clinician.

What the agent does

1. Reads clinical notes for an encounter.
2. Reviews assigned codes.
3. Investigates whether documentation supports codes — and whether higher-value codes could be supported with minor clarifications.
4. Drafts a “clinical documentation improvement” query to the physician.
5. Tracks query response and updates billing.

Why it sells

- **Every hospital CFO cares deeply.** Median hospital margins are razor-thin; CDI programs are high-ROI.
- **Massive market** — \$X billion spent on CDI services and software annually.
- **Accenture Health practice is strong.**
- **Outcome-based pricing works beautifully** — we take a percentage of recovered revenue.

Why feasibility is 7/10 (not higher)

HIPAA + EHR integration (Epic, Cerner) is non-trivial. But for a POC, synthetic data works fine.

Why sellability is maxed

Because ROI is measured in direct revenue recovery, this is the easiest sale on the list.

#10 — The Internal Knowledge Investigator (for Complex Enterprises)

Industry: Horizontal — any company >10,000 employees **Novelty:** 9/10 • **Feasibility:** 8/10 • **Sellability:** 8/10

The problem they don't know they have

Large enterprises have knowledge fragmented across SharePoint, Confluence, Notion, Slack, Google Drive, email, Jira, service desk tickets, meeting recordings, network drives. Enterprise search (Glean, Guru, Coveo) retrieves documents. **Nobody investigates across them to answer an actual question.**

Typical employee question: *“What’s our policy on promoting an employee whose manager is on PIP, in a region that’s in a hiring freeze, and the employee is on an H1B?”* — the answer exists across 6 documents, 2 Slack threads, and a 2023 email that nobody can find.

What the agent does

1. Understands the question (multi-condition queries).
2. Plans an investigation — which sources likely have what part of the answer.
3. Executes searches, reads results, identifies gaps.
4. Cross-references conflicting sources (which one is authoritative, which is stale).
5. Synthesizes with citations and flags uncertainties.

Why it sells

- **CIO budget for “enterprise AI search” is real and growing.**
- **Every employee is a user.** Scales enormously.
- **Measurable:** time-to-answer, ticket deflection, repeat-question reduction.
- **Differentiates from Glean** — which is retrieval, not investigation.

Why the novelty angle matters

The moment you show a demo where the agent reasons across 6 different sources to answer one question, with citations, you’ve destroyed every “enterprise search” competitor in the room.

Why feasibility is 8/10

RAG + source connectors + agentic loop. You have everything you need. Each client needs connector work, but Bedrock + LlamaIndex/similar covers the basics.

The Ranking Matrix (Reference Table)

#	POC	Novelty	Feasibility	Sellability	Primary Buyer	Industries
1	Controls & Compliance Evidence Agent	8	9	10	Chief Compliance Officer	All regulated
2	SecOps Tier-1 Triage Agent	7	9	10	CISO	Horizontal
3	Shadow IT / SaaS Sprawl Investigator	9	8	8	CIO / CISO	Horizontal
4	M&A Due Diligence Risk Agent	9	7	9	M&A partners	PE, IB, Corp Dev
5	Pricing Leakage Agent	8	9	9	CFO	B2B industries
6	Regulatory Change Impact Agent	9	8	9	Chief Risk Officer	Banking, Pharma
7	Silent Data Failure Detection Agent	10	8	8	CDO	Data-mature orgs
8	RFP Response Agent	7	9	9	Head of Proposals	B2B sales-led
9	Clinical Documentation Integrity Agent	8	7	10	Hospital CFO	Healthcare
10	Internal Knowledge Investigator	9	8	8	CIO	Horizontal, large

If I Had to Pick 3 to Build Next (After Your Telco POC)

Optimizing for maximum strategic value inside Accenture:

1. The Controls & Compliance Evidence Agent (#1)

Your architecture was literally designed for this. Fastest to build, highest sellability, universal applicability. This is the layup.

2. The SecOps Tier-1 Triage Agent (#2)

Horizontal reach across every industry practice. Massive buyer urgency. Positions you as the go-to agentic person for security deals across the firm.

3. The Pricing Leakage Agent (#5) or M&A Due Diligence Agent (#4)

Pick based on which practice you want relationships with — if Finance/CFO, pick #5. If Strategy/M&A, pick #4. Both have the same magic property: **self-funding ROI stories that make the sale almost automatic.**

That gives you a portfolio of **4 reference POCs** (telco + compliance + secops + pricing/M&A) covering the five biggest buyer personas in the enterprise (CIO, CISO, CCO, CFO, Business Ops). That's a position of real strength.

One Final Strategic Note

Notice what every POC on this list has in common:

They're all problems where a senior human currently does multi-step investigative work across scattered systems, and where audit-defensibility matters.

That's your architecture's sweet spot. **Don't chase POCs outside this pattern.** Creative writing agents, code generation agents, image agents — those are different games with different winners. Your edge is "investigative agent with citations and approval gates." Stay on that edge and you win.

Want me to go deep on the architecture + demo storyboard for any one of these? The compliance evidence agent (#1) is the one I'd pick first — and I can sketch what the 2-week build looks like.